

NoSQL Databases

Joshua Eckroth

<http://www.cse.ohio-state.edu/~eckroth/nosql-guide.pdf>

Outline

CAP theorem, ACID, BASE

Key-value database

Document database

Column-family store

Graph database

Object database

Choosing a DB

Trends

Resources

CAP theorem

From Eric Brewer, 2000.

- ▶ **Consistency**: all nodes see the same data at the same time
- ▶ **Availability**: every query returns a response
- ▶ **Partition tolerance**: continues to operate if nodes fail or message loss, or nodes are added/removed
- ▶ **Pick 2.**

E.g., if you want to scale up (A+P), must give up on consistency (C).

ACID

Regarding a transaction,

- ▶ **Atomicity**: all or nothing
- ▶ **Consistency**: transaction does not violate foreign key checks or other constraints
- ▶ **Isolation**: executing many transactions concurrently is same as serially; they don't interact
- ▶ **Durability**: completed transaction remains after power loss, etc. (i.e., written to disk)

ACID

- ▶ hard to achieve if transaction spans nodes
- ▶ hard to achieve without locking (which is detrimental to performance)
- ▶ give up C+I for availability, graceful degradation, and performance
- ▶ maybe even give up D for extra performance

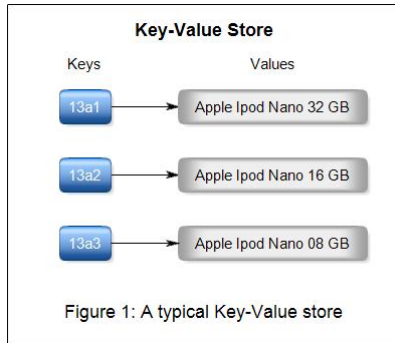
RDBMS's are typically ACID-compliant; NoSQL systems typically aren't.

BASE

- ▶ **Basically Available:** usually works
- ▶ **Soft state:** not always consistent
- ▶ **Eventually consistent:** reads across all nodes will eventually agree (if no updates happen in the meantime)

...as opposed to ACID. (Obviously a “backronym.”)

Key-value database



- ▶ can only retrieve by a unique key
- ▶ just get back a value; your client code interprets the value (db sees it as a blob)
- ▶ performance is uniform and fast
- ▶ horizontal scaling (seems to me) straightforward

Key-value database

Project Voldemort

Auto-replication, auto-partitioning, “tunable” consistency, transparent failure handling.

API:

- ▶ `get(key)` — returns a value
- ▶ `put(key, value)`
- ▶ `delete(key)`

<http://www.project-voldemort.com>

Key-value database

Redis

- ▶ master-slave replication; a slave can become a master if the master dies
- ▶ supports sets, lists, dictionaries

<http://redis.io/>

Memcached

- ▶ in-memory cache, never saved to disk
- ▶ when full, purges by LRU
- ▶ Facebook apparently has terabytes of “in-memory cache”

<http://memcached.org/>

Document database

Storage of arbitrary dictionaries that represent documents.

- ▶ also retrieved by key, or simple queries on fields
- ▶ data structures are stored (dictionaries & lists, each can be nested in the other)
- ▶ can store a list of keys, and retrieve those recursively, in one request
- ▶ good for syncing (copy newer revisions; propagate deletes)

Document database

Example “document”:

```
{
  '_id':      '29a8f708e',
  '_rev':     12,
  'author':   'Josh',
  'title':    'My first blog post',
  'tags':     ['foo', 'bar'],
  'content':  'Welcome to my blog! ...',
  'backlinks': ['37dd04387', '883bc2ccd']
}
```

Document database

MongoDB

- ▶ most popular NoSQL db
- ▶ easy sharding
 - ▶ sharding: data subsets stored in separate machines; not replicated; no joins; painful in RDBMS
- ▶ has some query support (find based on field values, plus some operators like $<$ $>$ etc.)
- ▶ can add indexes for faster queries

<http://www.mongodb.org/>

Document database

CouchDB

- ▶ create “views” of the data; these are updated when docs are updated
- ▶ if one doc is changed by two clients, two revisions are saved; merging is left to the client
- ▶ good for offline usage; changes are sync'ed later (again, no default merging)

<http://couchdb.apache.org/>

Column-family store

foo	username	email	gender
	foo	foo@bar.com	female
	614678718	24636234	5725652

bar	username	email	birthday	website
	bar	bar@foo.com	1970	foobar.com
	234523522	546625245	97198748	342534534

- ▶ key identifies a row in a table, which is part of 1+ column families
- ▶ each column family can have multiple columns
- ▶ values are timestamped (multiple versions of a value can be kept)

http://wikis.gm.fh-koeln.de/wiki_db/Datenbanken/Cassandra

Column-family store

Cassandra

Very high performance. Tunable consistency. Decentralized (no masters). No joins or subqueries.

API:

- ▶ `get(table, key, columnName)`
- ▶ `insert(table, key, rowMutation)`
- ▶ `delete(table, key, columnName)`

<http://cassandra.apache.org>

Graph database

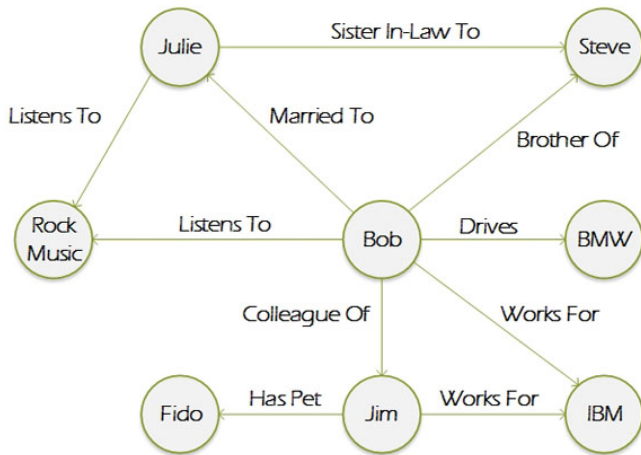
Nodes + edges that connect nodes.

- ▶ each node has arbitrary relations with others
- ▶ each node / edge has arbitrary properties
- ▶ can “walk” or query the graph according to these relations

Implementations:

- ▶ **Neo4J**: <http://www.neo4j.org>
- ▶ **HyperGraphDB**: <http://www.hypergraphdb.org>

Graph database



<http://www.computerweekly.com/feature/Whiteboard-it-the-power-of-graph-databases>

Object database

```
Foo f = new Foo();  
db.persist(f);  
for (Foo g : db.query("SELECT g FROM Foo")) {  
    // do something with g...  
}
```

- ▶ essentially, persisting live objects
- ▶ basic query support, e.g., find objects of this class, etc.
- ▶ sometimes suffers from poor indexing, poor search, memory/disk fragmentation

Implementations:

- ▶ **db4o**: <http://www.db4o.com>
- ▶ **Caché**: <http://www.intersystems.com/cache>

Choosing a DB

Do you need ACID compliance?

RDBMS	Key-value	Document	Column-family	Graph	Object
+1	-1	-1	-1	+1	+1

Do you expect your schema to change often?

RDBMS	Key-value	Document	Column-family	Graph	Object
-1	+1	+1	-1		+1

Do you expect to store terabytes / petabytes of data?

RDBMS	Key-value	Document	Column-family	Graph	Object
-1		-1	+1	+1	

Choosing a DB

Do you require syncing from mobile devices?

RDBMS	Key-value	Document	Column-family	Graph	Object
-1		+1	-1	-1	-1

Do you require horizontal scaling?

RDBMS	Key-value	Document	Column-family	Graph	Object
-1	+1	+1	+1		-1

Do you require extreme performance?

RDBMS	Key-value	Document	Column-family	Graph	Object
-1	+1		+1		

Choosing a DB

Do you require queries for arbitrary relations among data?

RDBMS	Key-value	Document	Column-family	Graph	Object
+1	-1	-1	-1	+1	

Do you want the DB to take care of complex constraints?

RDBMS	Key-value	Document	Column-family	Graph	Object
+1	-1	-1	-1	-1	

Do you want to save internal object state?

RDBMS	Key-value	Document	Column-family	Graph	Object
+1					+1

[RDBMS with Object-relational mapping (ORM)]

Trends

Search trends

<http://www.google.com/trends/explore?q=nosql#q=nosql&geo=US&date=1%2F2009%2049m&cmpt=q>

Job trends

<http://www.indeed.com/jobtrends?q=sql%2C+nosql&l=>

Job trends growth

<http://www.indeed.com/jobtrends?q=sql%2C+nosql&l=&relative=1>

Resources

Cassandra vs MongoDB vs CouchDB vs Redis vs Riak vs HBase vs Couchbase vs Neo4j vs Hypertable vs ElasticSearch vs Accumulo vs VoltDB vs Scalaris comparison

by Kristof Kovacs

[http://kkovacs.eu/
cassandra-vs-mongodb-vs-couchdb-vs-redis/](http://kkovacs.eu/cassandra-vs-mongodb-vs-couchdb-vs-redis/)

NoSQL Databases, a free book by Christof Strauch

<http://www.christof-strauch.de/nosql dbs.pdf>

The NoSQL Ecosystem, free book chapter by Adam Marcus

<http://www.aosabook.org/en/nosql.html>